



I am a senior R data engineer, I want to use a R program to send emails, which package you recommend me to use, and which email provider you advise to me.

Gemini 3.5 Thinking

As a senior R data engineer, your priorities are likely pipeline reliability, secure credential management, clean HTML rendering (for data reports), and minimal dependency overhead (especially avoiding Java issues in Docker/CI/CD pipelines).

Here are the top-tier R packages and email providers tailored for data engineering workflows.

Top Recommended R Packages

1. **blastula** (Highly Recommended)

Maintained by Posit, **blastula** is the modern industry standard for sending emails from R. It excels at creating beautiful, responsive HTML emails directly from R Markdown or Quarto documents.

- **Why it's great:** It allows you to embed `ggplot2` objects, tables, and markdown prose seamlessly. It also integrates perfectly with standard automation tools like Posit Connect, GitHub Actions, or Airflow.
- **Security:** It has built-in integration with the `keyring` package, making it easy to store and retrieve SMTP credentials securely without hardcoding them.
- **Usage Example:**

```
r

library(blastula)

# Compose the email
email <- compose_email(
  body = md("The daily **data pipeline report** is complete. See attached metrics."),
  footer = md("Sent via automated R script.")
)

# Send via SMTP
smtp_send(
  email,
  from = "pipeline@yourdomain.com",
```

```

to = "team@yourdomain.com",
subject = "Data Sync Status",
credentials = creds_smtp(
  host = "smtp.yourprovider.com",
  port = 465,
  user = "username",
  use_ssl = TRUE
)
)

```

2. emayili (Best for Programmatic/MIME Building)

If you prefer a clean, piping-compatible (`|>`) syntax for programmatically building complex MIME messages with attachments, **emayili** is an excellent choice.

- **Why it's great:** It is lightweight, actively maintained, and handles attachments, inline images, and multi-part text/HTML payloads very predictably.

3. gmailr (Best for Google Workspace Ecosystem)

If your organization is strictly on Google Workspace and you want to avoid dealing with SMTP or App Passwords.

- **Why it's great:** It interacts directly with the Gmail REST API using OAuth2, making it highly secure and compliant with strict corporate IT policies.

 **A Note on Legacy Packages:** Avoid **mailR**. It relies heavily on **rJava**, which frequently introduces severe configuration headaches in Linux environments, Docker containers, and CI/CD runners.

Recommended Email Providers

For engineering pipelines, you should skip traditional inbox providers (like a standard Outlook or Gmail account) and opt for a dedicated **transactional email service**. They offer API access, high deliverability rates, and tracking metrics.

Provider	Best For	Pros	Cons
Amazon SES (Simple Email Service)	AWS-native stacks & cost efficiency	Unbeatable pricing (\$0.10 per 10k emails); highly scalable.	Slightly more complex initial setup (DKIM/SPF domain verification).
Twilio SendGrid	High-volume API throughput	Robust R/HTTP API documentation; clean dashboard for debugging bounces.	Can get expensive at massive volumes; strict compliance algorithms.

Provider	Best For	Pros	Cons
Postmark	Bulletproof deliverability	Industry-best delivery speeds and detailed transit logs for debugging.	Priced higher than competitors; strictly for transactional mail.
Mailgun	Multi-region & log retention	Excellent log visibility and built-in email validation.	UI can be cluttered; pricing model scales quickly.

Architectural Recommendation

- **If you are hosted in AWS:** Use **Amazon SES** via SMTP credentials passed to `blastula` . It keeps your network latency minimal and your billing consolidated.
- **If you want the quickest setup with great logs:** Go with **SendGrid** or **Postmark**. You can configure a verified sending domain in 15 minutes, get an API/SMTP key, and monitor your pipeline's outbound delivery immediately.

What kind of payloads are you planning to send out—are they basic text notifications for pipeline failures, or are you looking to embed rich HTML data reports and attachments?

2026-06-20 10:49:54

You Asked



I only send emails to my own emails address. I want to use free email providers, I have gmail, outlook, and QQ emails. advise me.



Gemini 3.5 Thinking

Since you are sending emails exclusively to yourself for pipeline alerts, logs, or reports, you don't need a dedicated transactional provider. Your existing personal accounts are perfect for this, and deliverability won't be an issue because the receiving server automatically trusts emails authenticated from its own user.

Here is how your three current options stack up for automated R scripting, along with the best way to implement them.

1. Gmail + gmailr (Best Overall & Most Secure)

Google has strictly deprecated basic password authentication for SMTP, meaning you cannot just plug your regular Gmail password into a script. However, Gmail is still the best option if you leverage its API.

- **The Method:** Use the `gmailr` package.
- **How it works:** You create a free project in the Google Cloud Console, enable the Gmail API, and download a credentials JSON file. `gmailr` uses OAuth2 to authenticate.
- **Why it's great for an engineer:** Once the initial OAuth token is cached, it is incredibly stable. It completely bypasses SMTP ports, which are frequently blocked by cloud providers (like AWS EC2 or DigitalOcean) or corporate firewalls.

2. QQ Mail + `blastula` (Easiest SMTP Setup)

If you want a traditional SMTP setup without dealing with Google Cloud API configurations, QQ Mail is surprisingly developer-friendly for personal scripts.

- **The Method:** Use `blastula` or `emayili`.
- **How it works:** You log into the QQ Mail web interface, go to *Settings -> Account*, and generate an **Authorization Code (授权码)** for POP3/SMTP. You use this 16-character code as your password in R.
- **Why it's great:** It works out of the box with standard SMTP code, and QQ Mail rarely flags self-to-self automated traffic as spam.
- **Configuration:**
 - **SMTP Server:** `smtp.qq.com`
 - **Port:** 465 (SSL) or 587 (TLS)

3. Outlook + `blastula` (The Middle Ground)

Like Gmail, Microsoft is aggressively phasing out basic SMTP authentication for personal accounts (`@outlook.com` , `@hotmail.com`), forcing modern authentication (OAuth2).

- **The Method:** If your account still allows it, you must turn on Two-Factor Authentication (2FA) in your Microsoft Security settings and generate an **App Password**.
- **Configuration:** * **SMTP Server:** `smtp-mail.outlook.com`
 - **Port:** 587 (TLS)
- **Downside:** Microsoft's automated risk-detection algorithms are notoriously aggressive. If your R script suddenly fires an email from a new cloud IP address, Microsoft will often lock the account or block the connection until you log in via a browser to verify your identity.

Summary Recommendation

- **Go with Gmail (`gmailr`)** if your script runs on a cloud server or CI/CD pipeline, as API calls are immune to the port-blocking and IP-verification headaches that plague SMTP.

- **Go with QQ Mail (blastula)** if you are running scripts locally or want the fastest, simplest SMTP setup without setting up a developer console.

Regardless of which you choose, make sure to use environment variables (`Sys.getenv()`) or the `keyring` package so you don't commit your App Passwords or API secrets to your code repository.

Where is this R script going to run—is it executing locally on your machine, or are you deploying it to a cloud environment?

2026-06-20 10:53:10 You Asked 

I would like to push my pipeline to Github Actions, and use automations to send me emails. Based on your replies, should I use gmail and gmailr? advise.

 Gemini 3.5 Thinking

For a personal pipeline running on **GitHub Actions (GHA)**, the recommendation shifts. **You should avoid gmailr and instead use blastula with an SMTP App Password (using either Gmail or QQ Mail).**

While `gmailr` and OAuth2 are highly secure for user-facing applications, they introduce significant friction in an automated, non-interactive CI/CD environment like GitHub Actions.

The architectural reasons for this choice, along with instructions for implementing it like a senior engineer, are detailed below.

Why `gmailr` is a Headache on GitHub Actions

1. **The Ephemeral Runner Problem:** GitHub Actions runs on fresh, short-lived virtual machines. Because `gmailr` relies on an interactive browser "OAuth dance" to cache a token on your local machine, you would have to manually export that token file, encode it to a string (like Base64), save it as a GitHub Secret, and write a script to decode and rebuild the file inside the runner on every single execution.
2. **The 7-Day Expiration Trap:** If you use a personal `@gmail.com` account, your Google Cloud Console project will default to a "Testing" status. For unverified personal apps, **Google expires OAuth refresh tokens every 7 days**. This means your GitHub Action will break every week, forcing you to re-authenticate locally and re-upload a new token secret to GitHub.

Why blastula + SMTP App Password is Superior for GHA

Using standard SMTP with blastula bypasses the token lifecycle entirely. By generating a static **App Password** (Gmail) or **Authorization Code** (QQ Mail), you get a credential that:

- Natively integrates with GitHub Secrets as a simple text string.
- Never expires unless you manually revoke it.
- Can be injected instantly into the runner's environment variables.

The Engineering Setup

The standard implementation for this workflow utilizes Gmail as the engine, though the exact same logic applies to QQ Mail (just change the host to smtp.qq.com and port to 465).

Step 1: Generate your App Password

1. Go to your Google Account settings -> Security.
2. Ensure **2-Step Verification** is turned on.
3. Search for **App Passwords** in the search bar.
4. Generate a password for "Other (Custom Name)" and call it something like GitHub Actions Pipeline . Copy the 16-character code.

Step 2: Configure GitHub Secrets

In your GitHub Repository, navigate to **Settings -> Secrets and variables -> Actions -> New repository secret**.

- **Name:** EMAIL_PASSWORD
- **Secret:** Paste the 16-character code here (no spaces)

Step 3: Write your R Script (pipeline.R)

blastula has a built-in function called creds_envvar() specifically designed for CI/CD environments. It looks for a specified environment variable at runtime, meaning you never risk hardcoding secrets.

```
r

library(blastula)

# 1. Run your data engineering pipeline logic here
# df <- run_pipeline()

# 2. Compose your automated report
email <- compose_email(
  body = md("### 🚀 Pipeline Execution Successful")
)
```

```

    The daily data synchronization job has completed without errors."),
    footer = md("Sent automatically via GitHub Actions.")
)

# 3. Send using environment variables
smtp_send(
  email = email,
  from = "your_email@gmail.com",
  to = "your_email@gmail.com", # Sending to yourself
  subject = "✅ GHA Alert: Pipeline Success",
  credentials = creds_envvar(
    user = "your_email@gmail.com",
    pass_envvar = "EMAIL_PASSWORD", # Tell blastula to look at Sys.getenv("EMAIL_PASSWORD")
    host = "smtp.gmail.com",
    port = 465,
    use_ssl = TRUE
  )
)

```

Step 4: Create your GitHub Actions YAML (`.github/workflows/pipeline.yml`)

In your workflow file, you map the GitHub Repository Secret to the step's environment variables (`env`) so the R script can see it.

```

yaml

name: Run Data Pipeline

on:
  schedule:
    - cron: '0 6 * * *' # Runs daily at 6:00 AM UTC
  workflow_dispatch: # Allows manual triggering

jobs:
  run-pipeline:
    runs-on: ubuntu-latest
    steps:
      - name: Check out repository
        uses: actions/checkout@v4

      - name: Set up R
        uses: r-lib/actions/setup-r@v2

      - name: Install Dependencies
        run: |
          Rscript -e 'install.packages(c("blastula", "rmarkdown"))'

      - name: Execute Pipeline and Email

```

```
env:
  EMAIL_PASSWORD: ${ secrets.EMAIL_PASSWORD } # Injects the secret safely
run: Rscript pipeline.R
```

Summary Advice

Stick to **Gmail + blastula + creds_envvar()** . It provides a production-grade setup that requires zero maintenance once deployed, keeps your code repository perfectly clean of credentials, and avoids the token expiration pitfalls of the Google API.

2026-06-20 11:32:06 You Asked 

I tried to run your example, but got the below error. I used ``pass_envvar = Sys.getenv("EMAIL_PASSWORD")``, and I have added EMAIL_PASSWORD into my local computer's environment variable. When I call ``Sys.getenv("EMAIL_PASSWORD")``, I can see the password is available.

```
<error>
The environment variable defined by `pass_envvar` doesn't exist.
</error>
```

Gemini 3.5 Thinking

Ah, the classic API misdirection! You stumbled into a very common trap with how **blastula** designs its arguments.

The issue is that the **pass_envvar** argument expects the **name (the key)** of the environment variable as a character string, *not* the actual evaluated value.

What Happened Under the Hood

When you wrote this:

```
r

pass_envvar = Sys.getenv("EMAIL_PASSWORD")
```

R evaluated `Sys.getenv("EMAIL_PASSWORD")` first and replaced it with your actual 16-character password (let's pretend it's `"abcd1234efgh5678"`).

So **blastula** received:


```
r
```

```
pass_envvar = "abcd1234efgh5678"
```

`blastula` then looked inside your system for an environment variable *named* `"abcd1234efgh5678"` . It couldn't find one, so it threw the error: *"The environment variable defined by `pass_envvar` doesn't exist."*

The Fix

You just need to pass the literal string name of the variable. `blastula` will handle calling `Sys.getenv()` internally.

Change your configuration back to this:

```
r
```

```
credentials = creds_envvar(  
  user = "your_email@gmail.com",  
  pass_envvar = "EMAIL_PASSWORD", # <--- Just the name of the key in quotes  
  host = "smtp.gmail.com",  
  port = 465,  
  use_ssl = TRUE  
)
```

Since your local `Sys.getenv("EMAIL_PASSWORD")` already returns the password correctly, making this change will allow the script to run perfectly both on your local machine and when it transitions to GitHub Actions.

2026-06-20 11:47:17 You Asked 

I want to learn more details about the YAML file, especially on schedule, how to run script every 30 minutes, besides, how to cache dependencies instead of install them all every time the script runs.

 Gemini 3.5 Thinking

Optimizing workflows for speed and cost is a core part of moving a pipeline from a local machine into a production environment.

To run a pipeline every 30 minutes and implement dependency caching, the GitHub Actions YAML must be modified.

1. Triggering an Action Every 30 Minutes

GitHub Actions utilizes standard POSIX cron syntax to handle scheduling. To trigger a job every 30 minutes, configure the `schedule` block as follows:

```
yaml

on:
  schedule:
    - cron: '*/*30 * * * *'
```

Syntax Breakdown

Field	Position	Value	Meaning
Minute	1st	*/*30	Step interval: Every 30th minute (0, 30)
Hour	2nd	*	Wildcard: Every hour
Day of Month	3rd	*	Wildcard: Every day of the month
Month	4th	*	Wildcard: Every month
Day of Week	5th	*	Wildcard: Every day of the week

⚠ Critical Production Caveats for GHA Schedules:

- The Queue Delay:** GitHub Actions executes cron schedules on a resource-availability basis. Runs are rarely exact; a job scheduled for `:30` might not spin up until `:34` or `:41` depending on current global GitHub server traffic. If absolute real-time strictness is required, GHA is the wrong tool.
- The 60-Day Dormancy Rule:** If a repository receives zero code commits for **60 consecutive days**, GitHub automatically disables all scheduled cron jobs to save resources.

2. High-Efficiency Dependency Caching in R

Using raw `install.packages()` inside an automated script is a major CI/CD anti-pattern. GitHub-hosted runners are clean Ubuntu environments; running `install.packages()` means R will attempt to compile packages and their underlying system libraries from source on every single run, inflating a 10-second script into a 5-minute resource drain.

The **r-lib/actions** ecosystem provides two robust ways to cache R packages natively without needing manual `actions/cache` blocks.

Approach A: The `setup-r-dependencies` Action (Simplest)

If a pipeline relies on a handful of standard CRAN packages, use `r-lib/actions/setup-r-dependencies@v2`. Under the hood, this action uses `pak` to map out dependencies, pulls pre-compiled binaries via the Posit Public Package Manager (RSPM), and **automatically hooks into GitHub's caching layer** without any manual configuration.

yaml

```
- name: Set up R
  uses: r-lib/actions/setup-r@v2
  with:
    use-public-rspm: true # Injects fast pre-compiled binaries for Linux

- name: Install and Cache Dependencies
  uses: r-lib/actions/setup-r-dependencies@v2
  with:
    packages: |
      any::blastula
      any::dplyr
```

On the first run, it will install the packages and save the cache. On the second run (30 minutes later), it checks if your package list has changed; if not, it restores the environment instantly.

Approach B: The `renv` Action (Best for Enterprise Reproducibility)

For professional data engineering, tracking dependencies via an `renv.lock` file is the industry standard. If the project directory contains an `renv.lock` file, use `r-lib/actions/setup-renv@v2`. It handles all local project directory configurations, system headers, and cross-run caching seamlessly.

yaml

```
- name: Set up R
  uses: r-lib/actions/setup-r@v2

- name: Setup renv and Restore Cache
  uses: r-lib/actions/setup-renv@v2
```

Comprehensive Production Blueprint

The following production-ready YAML combines a 30-minute interval, manual workflow triggers (for ad-hoc debugging), fast binary mapping, and fully automated dependency caching:

yaml

```
name: 30-Min Automated Data Pipeline

on:
  schedule:
    - cron: '*/*/* * * *' # Executes every 30 minutes
  workflow_dispatch:      # Retains the ability to click "Run workflow" manually in the UI

jobs:
  execute-pipeline:
    runs-on: ubuntu-latest

    steps:
      - name: Check out repository
        uses: actions/checkout@v4

      - name: Set up R Environment
        uses: r-lib/actions/setup-r@v2
        with:
          use-public-rspm: true # Forces fast Linux binaries instead of source compilation


      - name: Load Cached R Packages
        uses: r-lib/actions/setup-r-dependencies@v2
        with:
          packages: |
            any::blastula
            any::jsonlite

      - name: Run Pipeline Script
        env:
          EMAIL_PASSWORD: ${ secrets.EMAIL_PASSWORD }
        run: Rscript pipeline.R
```

2026-06-20 11:57:07

You Asked





I need your to review my current pipeline folder structure and file. tell me if it is correct or not. advise my the best practices according to your senior DE experiences.

see my structure in this attached screenshot.



Your structure looks highly professional. You have opted for a **"Packaged Pipeline" architecture** (indicated by the presence of `DESCRIPTION` , `.Rbuildignore` , and the `R/` folder) combined with a traditional script runner layout (`scripts/` , `data/` , `output/`).

For a senior R data engineer, this is an excellent design paradigm. It allows you to write modular, unit-testable functions inside the `R/` directory and execute them sequentially via a lightweight wrapper script inside `scripts/` .

Here is a detailed peer review of your layout, along with best practices to ensure it runs seamlessly in GitHub Actions.

🏗️ The Structural Review

1. The Brilliant Moves (Senior Approvals)

- **The `DESCRIPTION` File:** This is a massive win for your CI/CD pipeline. Because you have a `DESCRIPTION` file, the `r-lib/actions/setup-r-dependencies@v2` action we discussed earlier will **automatically scan your Imports: field** to identify and cache your dependencies. You don't have to maintain a hardcoded list of packages in your YAML file.
- **Separation of Concerns:** Dividing your pipeline into functional logic (`R/`) and runtime execution (`scripts/`) keeps your codebase highly maintainable.

2. Architectural Risks to Address

Because this project acts as both an R package and an operational data pipeline, you need to watch out for a few common friction points:

Director y/File	Risk Factor	Senior Engineering Best Practice
<code>data/</code> & <code>output/</code>	Large datasets bloating Git history or breaking package builds.	1. Ensure <code>data/*</code> and <code>output/*</code> are added to your <code>.gitignore</code> so raw production data is never pushed to GitHub. 2. Crucial: Add both folders to your <code>.Rbuildignore</code> file. If you ever run package checks or build tools, R will throw errors if it encounters unstructured data folders.
<code>scripts/</code>	Path resolution errors in CI/CD.	When running on GitHub Actions, the working directory defaults to the repository root. Ensure the R script execution step in your YAML explicitly targets this folder: <code>run: Rscript scripts/pipeline.R</code>
<code>R/</code>	Functions not being explicitly loaded.	Since your functions live in <code>R/</code> , your runner script inside <code>scripts/</code> should ideally load them using <code>pkgload::load_all()</code> or by sourcing them efficiently, rather than manually calling <code>source()</code> on every single file.

Recommended Code Adjustments

Update your `.Rbuildignore`

To prevent your data pipeline infrastructure from interfering with R's internal package building/documentation logic, ensure your `.Rbuildignore` looks like this:

```
text

^.*\.Rproj$
^\.Rproj\.user$
^data$
^output$
^scripts$
```

Your GHA Runner Script Setup (`scripts/pipeline.R`)

Inside your execution script, utilize the fact that this is a packaged project. Instead of writing multiple `source("R/fetch_data.R")` lines, leverage `pkgload` :

```
r
```

```
# scripts/pipeline.R

# 1. Load all functions inside the R/ directory seamlessly
pkgload::load_all(".")

# 2. Execute your clean pipeline functions
raw_data <- fetch_pipeline_data()
clean_data <- process_data(raw_data)

# 3. Write locally (ephemeral to the GHA runner)
saveRDS(clean_data, "output/cleaned_metrics.rds")

# 4. Trigger the blastula email we built earlier
send_pipeline_report()
```

Are you planning to store your historical pipeline datasets inside this same GitHub repository, or will your script pull from and write to an external database or cloud storage bucket (like AWS S3) during execution?

2026-06-20 12:03:18 You Asked 

In my YAML, I had this line. is it necessary?

<line>

with:

extra-packages: any::rcmdcheck

</line>

 Gemini 3.5 Thinking

Short answer: **No, it is almost certainly not necessary** for your specific workflow, and you should remove it.

Here is the engineering breakdown of why that line is there and why it is safe to delete.

What that line actually does

The `rcmdcheck` package is a tool used to execute R CMD check . This is the strict utility that checks an R package for CRAN compliance—verifying things like whether your documentation

matches your functions, checking for undocumented components, running formal unit tests, and ensuring there are no syntax anomalies across platforms.

When you see `extra-packages: any::rcmdcheck` in a GitHub Actions template, it is because that template was originally designed for **R Package Developers** who want to test their package every time they push code.

Why you should remove it for your Data Pipeline

1. **Massive Overhead for a 30-Minute Cron:** `rcmdcheck` has a massive tree of dependencies. Installing it forces your GitHub Actions runner to download and configure a dozen extra checking tools on every single run. Since your pipeline runs every 30 minutes, this adds pointless time and resource consumption to your cache building.
2. **You aren't shipping to CRAN:** You are building an operational data pipeline that happens to use a packaged directory structure for cleanliness. You do not need to run a full, rigorous CRAN check just to execute `Rscript scripts/pipeline.R`.
3. **It will likely fail anyway:** `R CMD check` is notoriously strict. Because your project contains non-standard package folders like `data/`, `output/`, and `scripts/`, `rcmdcheck` will throw warnings or errors about an ungrounded package structure unless you have perfectly calibrated your `.Rbuildignore`.

When *would* it be necessary?

The **only** reason to keep it is if you have a later step in your YAML file that looks like this:

```
yaml
- name: Check Package
  uses: r-lib/actions/check-r-package@v2
```

If your workflow does not contain that step, strip `any::rcmdcheck` completely out of your `packages` or `extra-packages` block. Your runner will spin up faster, consume less bandwidth, and focus entirely on what matters: executing your pipeline and sending your email.

2026-06-20 12:11:16 You Asked 

setup r environment took 32 seconds, the install dependencies 27s.
DOes this sound ok to you?

For an R-based pipeline running on GitHub Actions, **these numbers are excellent**. In fact, they are a massive green flag that your environment optimizations are working perfectly.

To give you the senior engineer's perspective on why these numbers are great, let's break down what is happening under the hood.

Setup R Environment: 32 Seconds (Standard Overhead)

This step is completely dependent on GitHub's infrastructure, not your code.

- Inside those 32 seconds, GitHub is spinning up a fresh Ubuntu virtual machine, downloading the specified version of R, installing system-level prerequisites, and configuring environment paths.
- Anything around **30 to 45 seconds** is the baseline standard for provisioning a clean R runtime on GHA. You cannot optimize this down much further, and you don't need to.

Install Dependencies: 27 Seconds (The Real Victory)

This is where your layout shines. In the R ecosystem, dependency installation is notorious for destroying CI/CD performance.

- If your workflow were downloading raw source code from CRAN and compiling packages (like `blastula`, `jsonlite`, `stringr`, etc.) on the fly, this step would easily take **3 to 7 minutes** because of C/C++ compilation.
 - A clock time of **27 seconds** is concrete proof that either your GitHub Action cache hit successfully, or the `use-public-rspm: true` flag is doing its job by pulling pre-compiled Linux binaries instantly.
-

The 30-Minute Interval Perspective

Since your pipeline is scheduled to run every 30 minutes, your total framework overhead before your actual script even starts executing is just under **1 minute**.

GitHub provides **2,000 free runner minutes per month** for public repositories (and standard pricing for private repos is highly affordable). At ~1 minute of overhead per run, 48 runs a day will consume roughly 1,400 to 1,500 minutes a month. You are well within a highly efficient, cost-effective operating window.

How long does the actual data processing and email sending step (`Rscript scripts/pipeline.R`) take to run once the environment is ready?



The schedule did not work as expected. I did not receive any emails after 30 minutes. besides, I did not see any execution log in the workflow run.

```
<schedule>
name: Run R Script

# on: [push, pull_request]
on:
  schedule:
    - cron: '*/*/* * * * *' # Executes every 30 minutes
  workflow_dispatch:      # Retains the ability to click "Run workflow"
                             manually in the UI

</schedule>
```

Gemini 3.5 Thinking

Welcome to one of the most notoriously frustrating quirks of GitHub Actions. When a schedule fails to run and leaves zero execution logs, it almost always means GitHub's backend scheduler hasn't registered or parsed your file yet.

Because your YAML syntax itself is completely correct, the issue boils down to how and where the file is saved.

Run through this senior engineering checklist to unblock the pipeline.

The "Silent Failure" Checklist

1. Is the file on your Default Branch? (Most Common Trap)

GitHub Actions **strictly ignores** the `schedule` trigger if the YAML file is sitting on a feature branch or development branch.

- **The Rule:** The workflow file must be committed and merged into your repository's **default branch** (usually `main` or `master`) for the cron timer to activate.
- **Fix:** Merge your pipeline branch into `main`.

2. Is the Folder Path Exactly Correct?

The path to your workflow file is highly case-sensitive and rigid. If there is a single typo in the hidden directory structure, GitHub will ignore the file entirely.

- **The Rule:** It must be exactly `.github/workflows/your-file-name.yml` .
- **Common Mistakes:** Naming the folder `.github/workflow/` (singular) or putting it in the root directory. Looking back at your folder screenshot, make sure inside `.github/` you have a `workflows/` subfolder.

3. Use `workflow_dispatch` to Force a Diagnostic Test

Since you wisely included `workflow_dispatch:` , you can completely bypass the cron scheduler to see if GitHub recognizes the file.

1. Go to your repository on GitHub.com.
2. Click on the **Actions** tab.
3. Look at the left sidebar. Do you see your workflow name ("**Run R Script**") listed there?

- **Scenario A: It is NOT listed.** GitHub hasn't parsed your file at all. This confirms **Check #1** or **Check #2** is the issue.
- **Scenario B: It IS listed.** Click on it, click the "**Run workflow**" dropdown button on the right, and trigger it manually. If it runs and sends you an email, your code is perfect, and the issue is simply **The GitHub Cron Delay** (detailed below).

The "GitHub Cron Delay" Reality Check

If your file is on `main` and works manually via `workflow_dispatch` , but the 30-minute timer hasn't fired, you are experiencing standard GitHub resource constraints.

GitHub explicitly states that the `schedule` trigger does not guarantee exact execution times. Because global cron demands spike exactly at the hour (`:00`) and half-hour (`:30`), GitHub queues these jobs.

- For a high-frequency cron like `*/30 * * * *` , it is incredibly common for the first run to take **45 to 70 minutes** to actually fire.
- Once the loop catches execution on the GitHub backend, it will stabilize, but it will almost never run at *exactly* 30-minute increments.

When you look at the **Actions** tab in GitHub, do you see the "Run R Script" workflow listed in the left sidebar?